

WHY I HATE CI

GRTECLYN UPDATE: PORTING THE MATTER CLASSES

✓ Update gpu-build.yml with CUDA 12.5

Browse files

🔗 training/202406_grtl_meeting

👤 julianakwan committed 5 days ago Verified

1 parent 82fdf08 commit d744e34

Showing 1 changed file with 1 addition and 1 deletion.

Whitespace Ignore whitespace Split Unified

2 .github/workflows/gpu-build.yml

@@ -29,7 +29,7 @@ jobs:	
29 AMREX_HOME: \${github.workspace}}/amrex	29 AMREX_HOME: \${github.workspace}}/amrex
30 BUILD_ARGS: USE_\${matrix.gpu-backend}}=TRUE DEBUG=\${matrix.debug}} USE_MPI=TRUE	30 BUILD_ARGS: USE_\${matrix.gpu-backend}}=TRUE DEBUG=\${matrix.debug}} USE_MPI=TRUE
31 CUDA_MAJOR_VERSION: 12	31 CUDA_MAJOR_VERSION: 12
32 - CUDA_MINOR_VERSION: 0	32 + CUDA_MINOR_VERSION: 5
33	33
34 steps:	34 steps:
35 - name: Checkout AMReX	35 - name: Checkout AMReX

✗ Merge branch 'training/202406_grchombo_meeting' of github.com:GRTLCol...

...laboration/GRTeclyn into training/202406_grchombo_meeting because I modified gpu-build.yml on the website interface

🔗 training/202406_grtl_meeting

👤 julianakwan committed 5 days ago

Showing 1 changed file with 1 addition and 1 deletion.

Whitespace Ignore whitespace Split Unified

2 .github/workflows/gpu-build.yml

@@ -29,7 +29,7 @@ jobs:	
29 AMREX_HOME: \${github.workspace}}/amrex	29 AMREX_HOME: \${github.workspace}}/amrex
30 BUILD_ARGS: USE_\${matrix.gpu-backend}}=TRUE DEBUG=\${matrix.debug}} USE_MPI=TRUE	30 BUILD_ARGS: USE_\${matrix.gpu-backend}}=TRUE DEBUG=\${matrix.debug}} USE_MPI=TRUE
31 CUDA_MAJOR_VERSION: 12	31 CUDA_MAJOR_VERSION: 12
32 - CUDA_MINOR_VERSION: 0	32 + CUDA_MINOR_VERSION: 5
33	33
34 steps:	34 steps:
35 - name: Checkout AMReX	35 - name: Checkout AMReX



HOW TO AMREXIFY EXISTING GRCHOMBO CODE

- ▶ BoxLoops -> ParallelFor
- ▶ Some functions need to be prefaced by macros so that CUDA/HIP knows they belong to the device or host.

```
72 // Calculate CCZ4 right hand side
73 if (m_p.max_spatial_derivative_order == 4)
74 {
75     BoxLoops::loop(CCZ4RHS<MovingPunctureGauge, FourthOrderDerivatives>(
76         m_p.ccz4_params, m_dx, m_p.sigma, m_p.formulation),
77         a_soln, a_rhs, EXCLUDE_GHOST_CELLS);
78 }
```

GRChombo

```
105 // Calculate CCZ4 right hand side
106 if (simParams().max_spatial_derivative_order == 4)
107 {
108     CCZ4RHS<MovingPunctureGauge, FourthOrderDerivatives> ccz4rhs(
109         simParams().ccz4_params, Geom().CellSize(0), simParams().sigma,
110         simParams().formulation);
111     amrex::ParallelFor(
112         a_rhs,
113         [=] AMREX_GPU_DEVICE(int box_no, int i, int j, int k) {
114             ccz4rhs.compute(i, j, k, rhs_arrs[box_no], soln_c_arrs[box_no]);
115         });
116 }
117 }
```

GRTEclyn

```
#define AMREX_GPU_HOST      __host__
#define AMREX_GPU_DEVICE   __device__
#define AMREX_GPU_GLOBAL   __global__
#define AMREX_GPU_HOST_DEVICE __host__ __device__
```


HOW TO AMREXIFY EXISTING GRCHOMBO CODE

- ▶ A lot of compute functions now have the form (int i, int j, int k, const amrex::Array4<data_t> &out, const amrex::Array4<data_t const> &in) instead of passing in a Cell object
 - ▶ inner const -> data is read only
 - ▶ outer const -> can't reposition pointer
- ▶ Member functions are now public so they are discoverable by CUDA (these were protected in GRChombo)

```
template <class matter_t, class gauge_t, class deriv_t>
template <class data_t>
void MatterCCZ4BHS<matter_t, gauge_t, deriv_t>::compute(
    Cell<data_t> current_cell) const
```

GRChombo

```
template <class matter_t, class gauge_t, class deriv_t>
template <class data_t>
AMREX_GPU_DEVICE AMREX_FORCE_INLINE void
MatterCCZ4BHS<matter_t, gauge_t, deriv_t>::compute(
    int i, int j, int k, const amrex::Array4<data_t> &rhs,
    const amrex::Array4<data_t const> &state) const
```

GRTEclyn

PORTING THE MATTER CLASSES


- ▶ MatterCCZ4RHS: usual CCZ4RHS but includes two extra variables for the matter field and conjugate momentum
- ▶ ScalarField class: matter_t object that is a specific type for many of these other objects
- ▶ DefaultPotential class: empty potential $V = 0$
- ▶ Potential: has a mass term also
- ▶ MatterWeyl4
- ▶ EMTensor
- ▶ MatterConstraints: includes Ham and Mom
- ▶ ScalarBubble: generates initial condition -> moved to InitialConditions/ScalarField directory
- ▶ FixedGridTagging : (incidental) the ScalarField example had fixed grids, so I ported this also

UNIT TESTING

- ▶ These classes have been tested under the doctest framework:
 - ▶ BSSNMatterTest
 - ▶ Run using `-dt-tc="Matter BSSN"`
 - ▶ Tests for:
 - ▶ RHS calculation
 - ▶ Hamiltonian and the momentum vectors ('Ham', 'Mom1', ...)
 - ▶ MatterWeyl4
 - ▶ Run using `-dt-tc="Matter Weyl4"`:
 - ▶ Tests for:
 - ▶ EMTensor ('rho')
 - ▶ Weyl4 calculation ('WeylRe' and 'WeylIm')
- ▶ Note that I have a slightly different approach to the HDF5 file dump to Miren - all the fields are named in the HDF5 file so you can separate them for debugging using `h5dump` or `h5ls` or `h5py`.

SCALAR FIELD EXAMPLE

- ▶ In GRChombo, this was a scalar bubble + Kerr BH -> only the scalar bubble remains
- ▶ Same functions you know and love:
 - ▶ `specificAdvance` -> `same`
 - ▶ `initialData` -> `initData`
 - ▶ `specificEvalRHS` -> `same`
 - ▶ `specificUpdateODE` -> `same`
 - ▶ `computeTaggingCriterion` -> `errorEst`
 - ▶ `specificPostTimeStep` ->
 - ▶ `*new*` derive
- ▶ To see member functions/inheritance diagrams, refer to the AMReX doxygen pages



THE GITHUB REPOSITORY

- ▶ [GitHub.com/GRTLCollaboration/GRTeclyn](https://github.com/GRTLCollaboration/GRTeclyn)
- ▶ 3 branches:
 - ▶ enhancement/matter_class:
 - ▶ updated matter classes, unit tests
 - ▶ enhancement/scalar_fields:
 - ▶ updated matter classes, no unit tests, and the scalar field example in the Examples/directory
 - ▶ training/202406_grtl_meeting: hands-on exercise for this afternoon

HOW TO RUN


- ▶ The readme on training/202406_grtl_meeting contains a comprehensive guide to running the scalar field example on Dawn.
- ▶ For other systems, I have found that the old GRChombo module files work well
 - ▶ e.g. on Icelake nodes, source ~/GRChombo/InstallNotes/MakeDefsLocalExamples/CSD3-Icelake-Intel-modules.sh, then set `USE_COMP=intel-llvm`
- ▶ Example/ScalarFields includes a sample params.txt file. But you must change the output directory! Also notice that some variables have been commented out because they use the default value.

REGRESSION TESTING


- ▶ How do you know if you have the right output?
- ▶ There is a `params_test.txt` in the directory and I have also uploaded a small output file to `.github/workflows/data`
- ▶ AMReX provides a tool called `fcompare`:
 - ▶ `fcompare.<comp>.ex plt_00003 .github/workflows/data/plt_scalar_field_compare_00003`
- ▶ For more details see: https://amrex-codes.github.io/amrex/docs_html/Post_Processing.html#fcompare

VIEWING THE OUTPUTS

- ▶ ParaView
 - ▶ plt* files; select AMReX/BoxLib grid reader
- ▶ yt
 - ▶ `import yt; my_data = yt.load("plt...")`
- ▶ The snapshot tool included with AMReX:
 - ▶ Go to `~/amrex/Tools/Plotfile`; build with `make`
 - ▶ Run with `./fsnapshot.<comp>.ex -v phi -p Palette plt_00003`
 - ▶ The output is in the parent directory of plt...
 - ▶ https://amrex-codes.github.io/amrex/docs_html/Post_Processing.html#fsnapshot



ParaView



fsnapshot

FUTURE WORK

- ▶ CI fixes: the gpu-build workflow fails (although it will compile with CUDA and SYCL on CSD3)
- ▶ Parameter file could be cleaned up e.g. params.txt has some unused options, plot_vars is currently ignored
- ▶ For historical reasons, some files are called NewMatterConstraints etc. and these need to be cleaned up.
- ▶ Support for multiple scalar fields