

The OneAPI ecosystem: Intel Advisor, SYCLomatic and you”.

Stephen Blair-Chappell

Certified Intel oneAPI Instructor



About Me

Stephen Blair-Chappell



- Freelance software consultant
- Intel Certified oneAPI instructor.
- Formerly the Technical Director at Bayncore
- For 18 years he was a Technical Consulting Engineer at Intel
- Author of the book "Parallel Programming with Intel Parallel Studio XE".
- Sponsored by Lenovo to do some work at their Centre's of Excellence.

When I'm not working

- Trustee of two charities involved with mental health & young people.
- Adventurous DIY
- Play the Pipe Organ

oneAPI Spec Elements

The Spec is made of 7 core elements.



oneDPL

oneAPI Data Parallel C++ Library

A companion to the DPC++ Compiler for programming oneAPI devices with APIs from C++ standard library, Parallel STL, and extensions.



oneDNN

oneAPI Deep Neural Network Library

High performance implementations of primitives for deep learning frameworks.



oneCCL

oneAPI Collective Communications Library

Communication primitives for scaling deep learning frameworks across multiple devices.



Level Zero

oneAPI Level Zero

System interface for oneAPI languages and libraries.



oneDAL

oneAPI Data Analytics Library

Algorithms for accelerated data science.



oneTBB

oneAPI Threading Building Blocks

Library for adding thread-based parallelism to complex applications on multiprocessors.



oneMKL

oneAPI Math Kernel Library

High performance math routines for science, engineering, and financial applications.

<https://uxlfoundation.org/>

Intel® oneAPI Toolkits




Intel® oneAPI Base Toolkit



A core set of high-performance libraries and tools for building C++, SYCL, C/OpenMP, and Python applications

Add-on Domain-specific Toolkits




For HPC developers

Intel® oneAPI Tools for HPC
Deliver fast Fortran, OpenMP & MPI applications that scale



For visual creators, scientists & engineers

Intel® oneAPI Rendering Toolkit
Accelerate visual compute, deliver high-performance, high-fidelity visualization applications.



For edge & IoT developers

Intel® oneAPI Tools for IoT
Build efficient, reliable solutions that run at network's edge

Toolkits powered by oneAPI



For AI developers & data scientists

Intel® AI Analytics Toolkit
Accelerate machine learning & data science pipelines end-to-end with optimized DL & ML frameworks & high-performing Python libraries



For deep learning inference developers

Intel® OpenVINO™ toolkit
Deploy high performance inference & applications from edge to cloud

Download at intel.com/oneAPI
Or visit Intel® DevCloud for oneAPI

Intel® oneAPI Base Toolkit

Accelerate Data-centric Workloads

A core set of core tools and libraries for developing high-performance applications on Intel® CPUs, GPUs, and FPGAs.

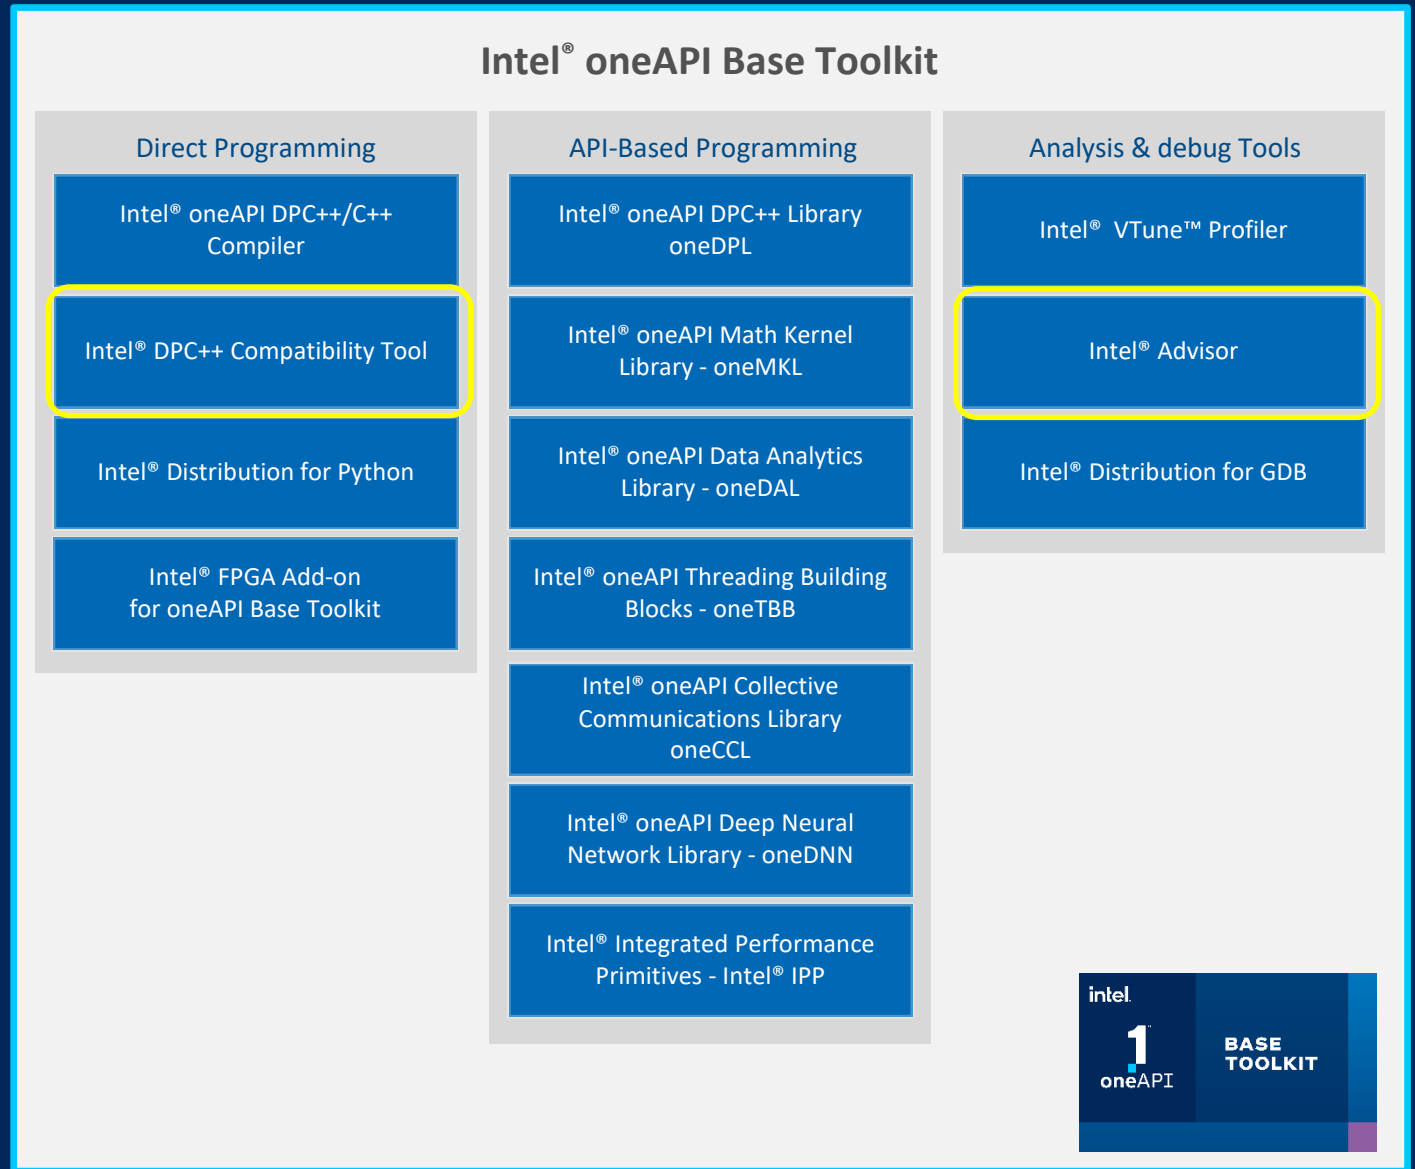
Who Uses It?

- A broad range of developers across industries
- Add-on toolkit users since this is the base for all toolkits

Top Features/Benefits

- Data Parallel C++ compiler, library and analysis tools
- SYCLomatic / DPC++ Compatibility tool helps migrate CUDA code to C++ with SYCL
- Python distribution includes accelerated scikit-learn, NumPy, SciPy libraries
- Optimized performance libraries for threading, math, data analytics, deep learning, and video/image/signal processing

[Learn More & Download](#)





Intel Advisor

|

Intel® Advisor

Configure Your Accelerated Computing Solution

Offload Advisor

Estimate performance of offloading to an accelerator

Roofline Analysis

Optimize CPU/GPU code for memory and compute

Vectorization Advisor

Add and optimize vectorization

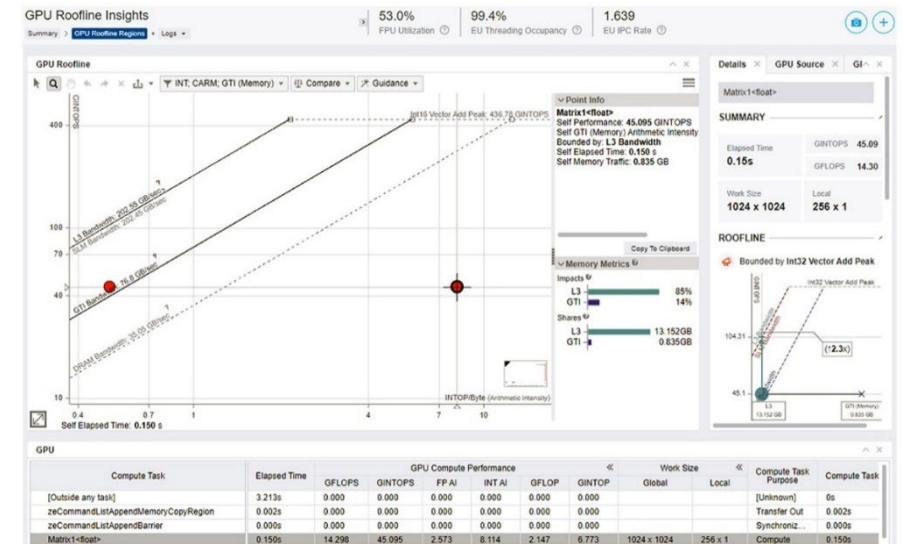
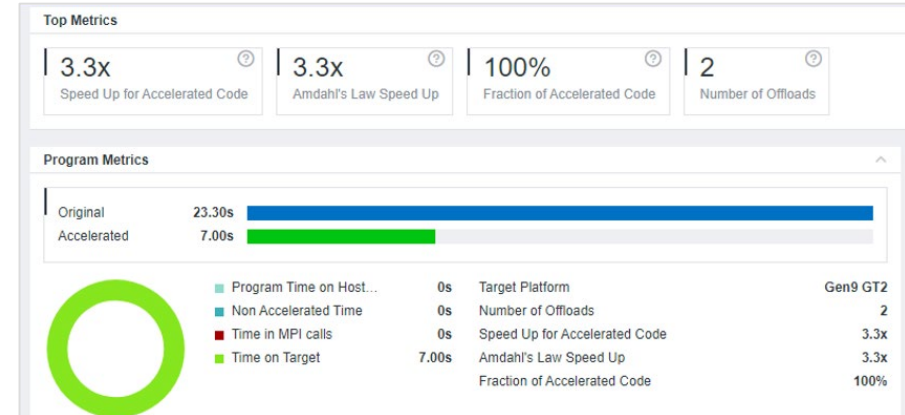
Threading Advisor

Add effective threading to unthreaded applications

Flow Graph Analyzer

Create and analyze efficient flow graphs

[Learn More & Download](#)

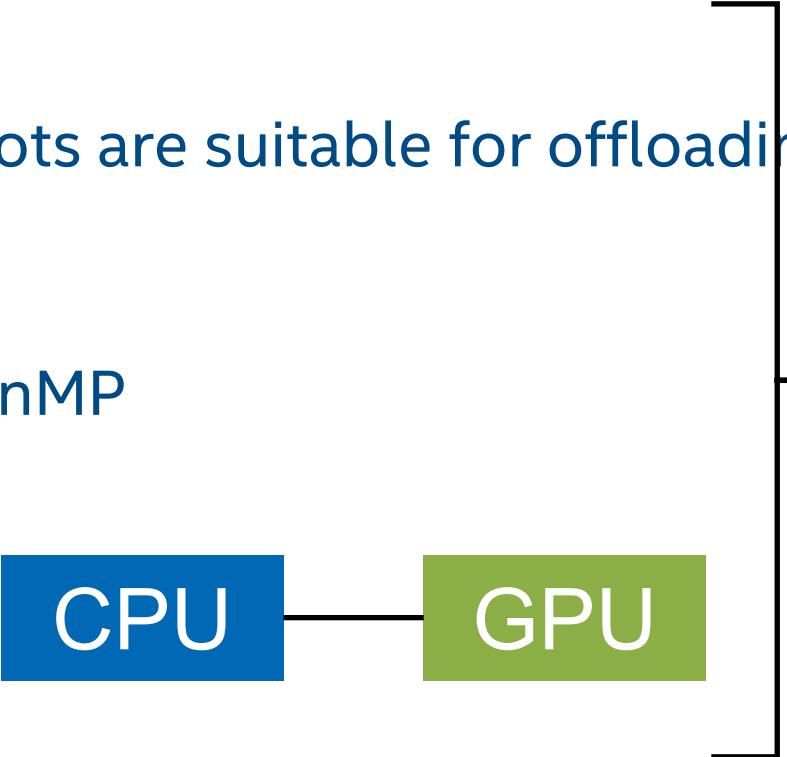


Know Before You Offload

Advisor Offload Modelling



Offload Modelling – doing it by hand

- 1 Run code on CPU and find hotspots **CPU**
 - 2 Examine results – decide which hotspots are suitable for offloading
 - 3 Implement offload using SYCL or OpenMP
 - 4 Run code on CPU and GPU and see if there is a speed up
- 
- The diagram illustrates the offload modelling process. It features a list of four steps. Step 1 is 'Run code on CPU and find hotspots' with a blue box labeled 'CPU' to its right. Step 2 is 'Examine results – decide which hotspots are suitable for offloading'. Step 3 is 'Implement offload using SYCL or OpenMP'. Step 4 is 'Run code on CPU and GPU and see if there is a speed up'. To the right of steps 2, 3, and 4, there is a vertical bracket. To the left of this bracket, there is a horizontal line connecting a blue box labeled 'CPU' to a green box labeled 'GPU'.

Potentially time-consuming.
Possibly no return-on-investment

Offload Modelling – with Advisor

1

Run code on CPU using Advisor
(multiple stage profiling collection)

2


Examine results – decide which hotspots are suitable for offloading

3

Implement offload using SYCL or OpenMP

4

Run code on CPU and GPU
and see if there is a speed up



Advisor
Simulates
all these
stages



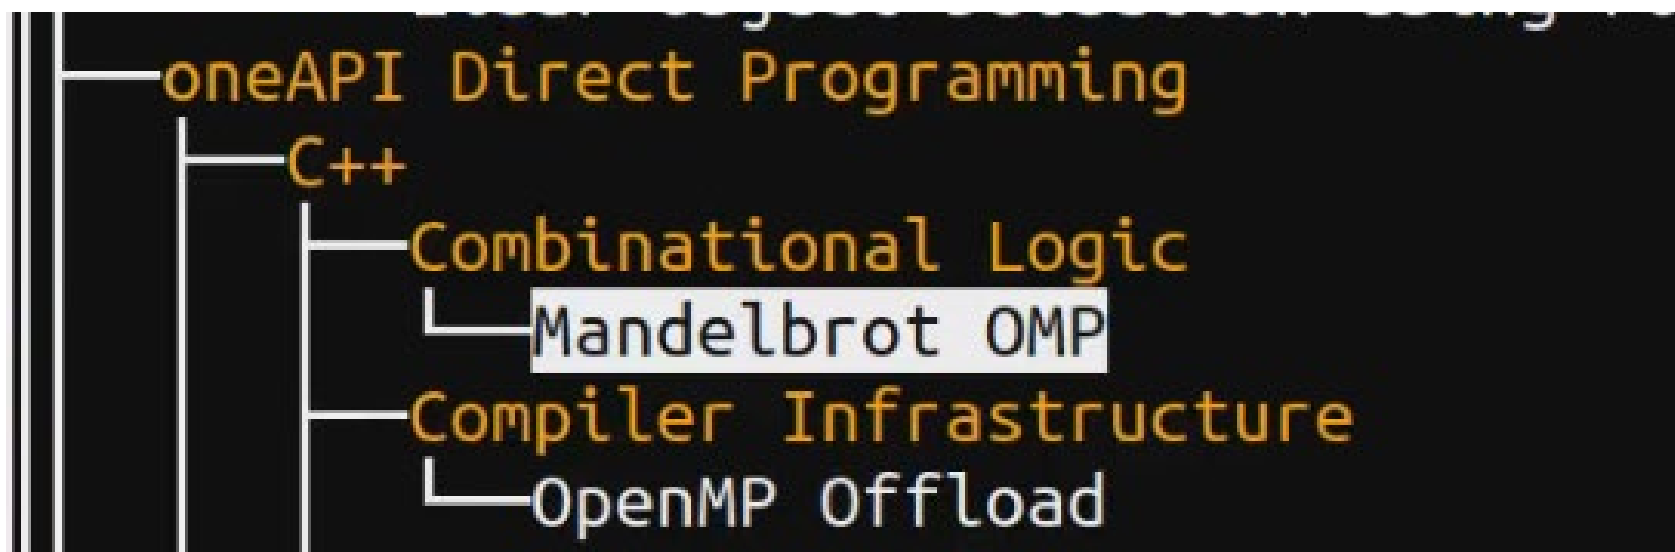
Report

CPU

CPU

Demo – An example you could experiment with

- `git clone https://github.com/oneapi-src/oneAPI-samples.git`




```
make CXX=icpx EXTRA_CFLAGS=-g
```

```
advisor --collect=offload -- ./release/Mandelbrot
```

OUR EXAMPLE

LAMMPS Molecular Dynamics Simulator



<https://www.lammps.org/>

- DO THE DEMO NOW!

Using accuracy presets to control modelling

- Default (Medium Accuracy)

```
advisor --collect=offload --config=gen12_tgl
--project-dir=./cpu2gpu_offload_modeling --
./release/Mandelbrot 1
```

- Low Accuracy

```
advisor --collect=offload -accuracy=low
--config=gen12_tgl --project-
dir=./cpu2gpu_offload_modeling --
./release/Mandelbrot 1
```

- Getting list of steps

```
advisor --collect=offload --dry-run
--config=gen12_tgl --project-
dir=./cpu2gpu_offload_modeling --
./release/Mandelbrot 1
```

| Low | Medium | High |
|--|--|---|
| 5-10x overhead | 15-50x overhead | 50-80x overhead |
| Survey Trip Count Offload Modelling | Survey Trip Count Offload Modelling | Survey Trip Count Dependency analysis Offload Modelling |
| L1 Cache | L1 Cache + Host- Device data | L1 Cache + Host- Device data |

Steps to Offload Projection with Advisor

1. Run a **Survey**: get a list of hotspots

```
advisor -collect survey ...
```

- Sampling
- Binary Static Analysis
- Compiler & debug info

2. Run a **Trip Count**: count loop iteration

```
advisor -collect=tripcounts -target-device=gen9_gt2 .
```

- Trip count
- Cache simulation

3. Perform a **dependency analysis** [optional for quick modelling]

```
advisor -collect dependencies . . .
```

- Check memory accesses
- Loop selection heuristic

4. **Model** the Performance

```
advisor -collect projection -no-assume-dependencies . . .
```

- Generate HTML report

Expensive Steps

GPU-GPU modelling

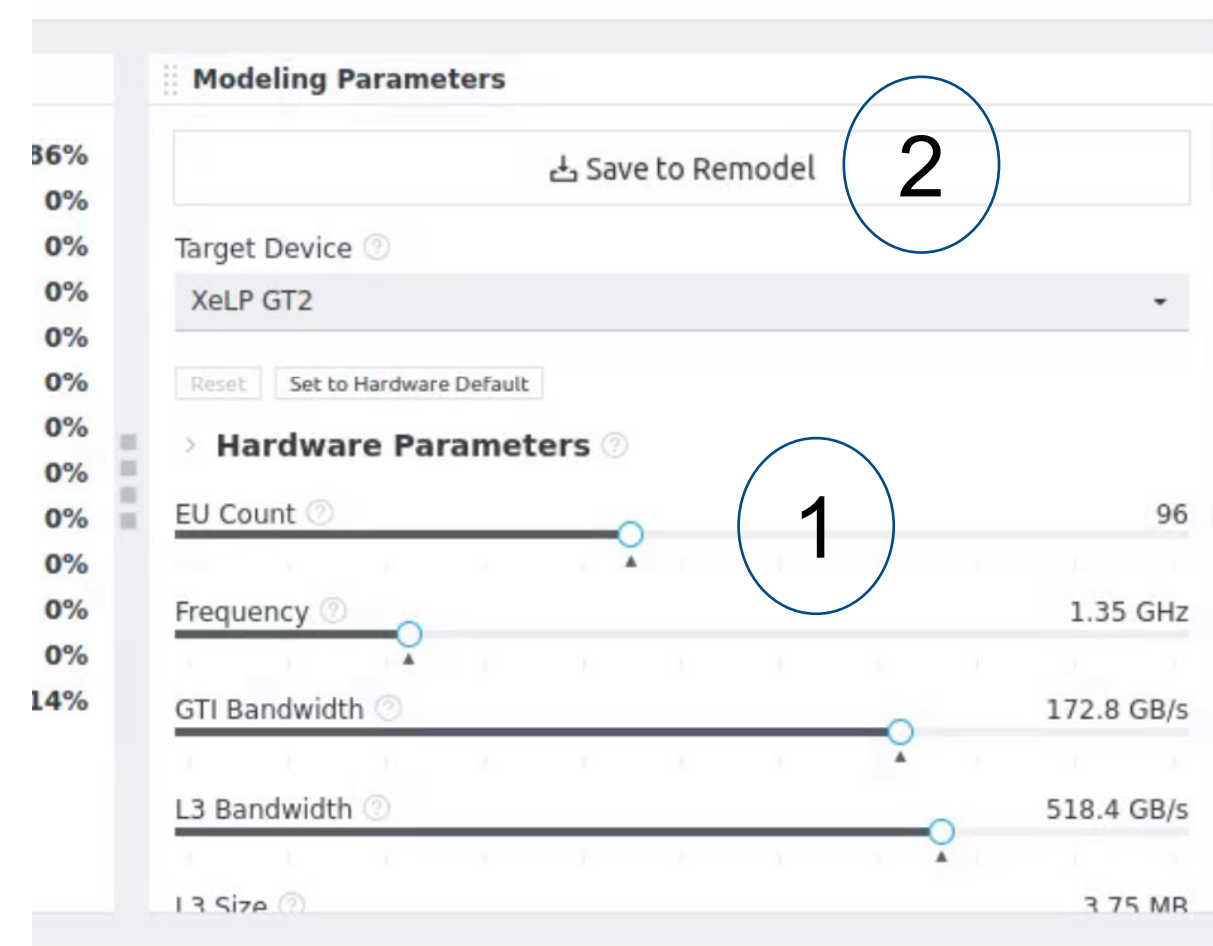
- Add the `-gpu` flag
- Runs code on ACTUAL GPU and models against new GPU.
- Use when upgrading from one GPU to another one.
- NB: use `-accuracy=low`

Comparison of CPU-GPU and GPU-GPU modelling

| CPU-GPU | GPU-GPU |
|------------------|--|
| Survey | Survey (on GPU) |
| Trip Count | Trip Count (Characterization - num Floats and Integer operations) |
| Dependency Check | x |
| Model Offloading | Model Offloading |

Re-modelling for a different GPU

- once you already have a set of results



1. Adjust Values
2. Save parameters
3. Re-run modelling

e.g.: `advisor -c=projection`
`--custom-`
`config=config.toml`
`--config=gen12_tgl`

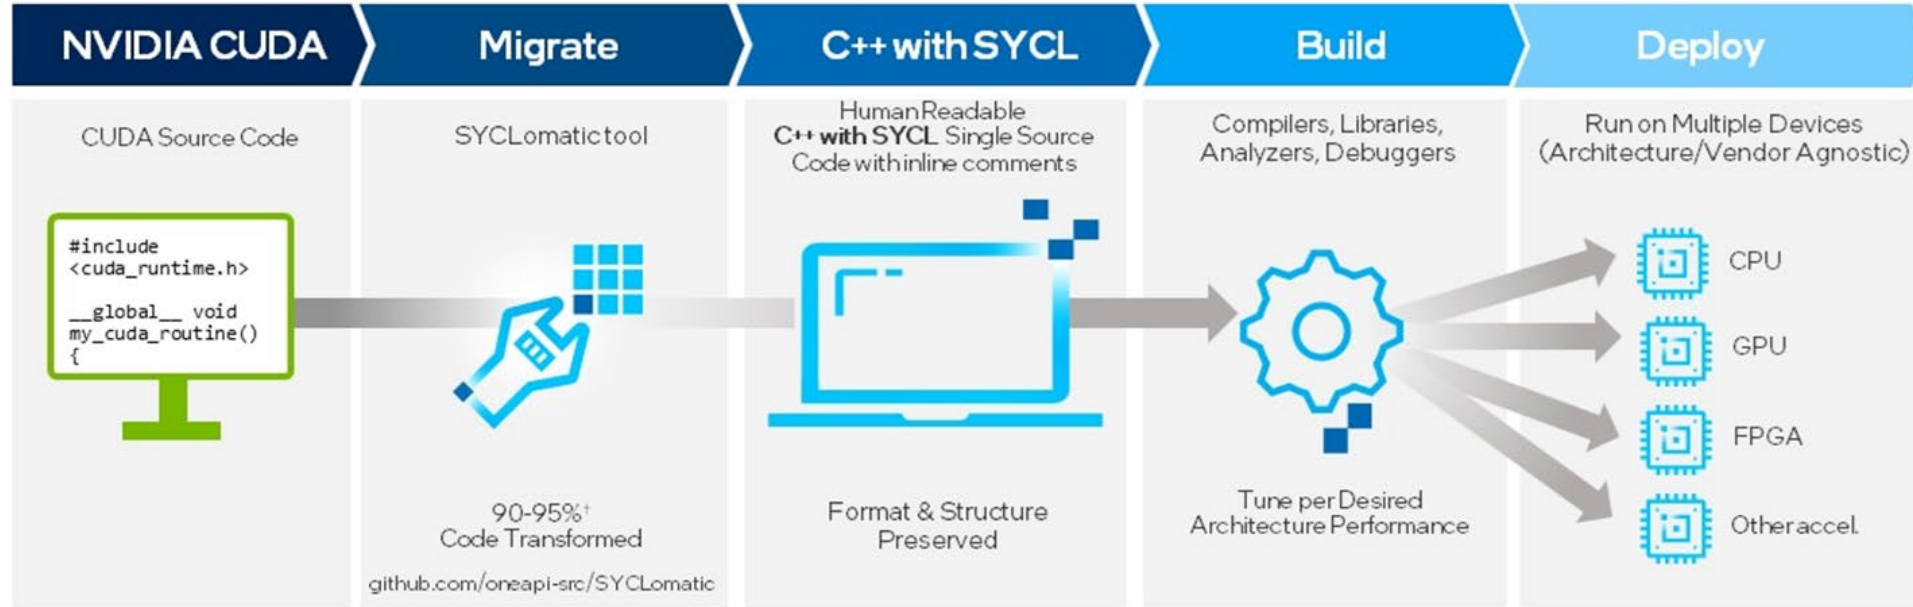
Advantage: quicker than doing a completely new modelling



SYCLomatic

CUDA to SYCL Migration Made Easy

Open Source SYCLomatic Tool Reduces Code Migration Time



Assists developers migrating code written in CUDA to C++ with SYCL, generating **human readable** code wherever possible

~90-95% of code typically migrates automatically¹

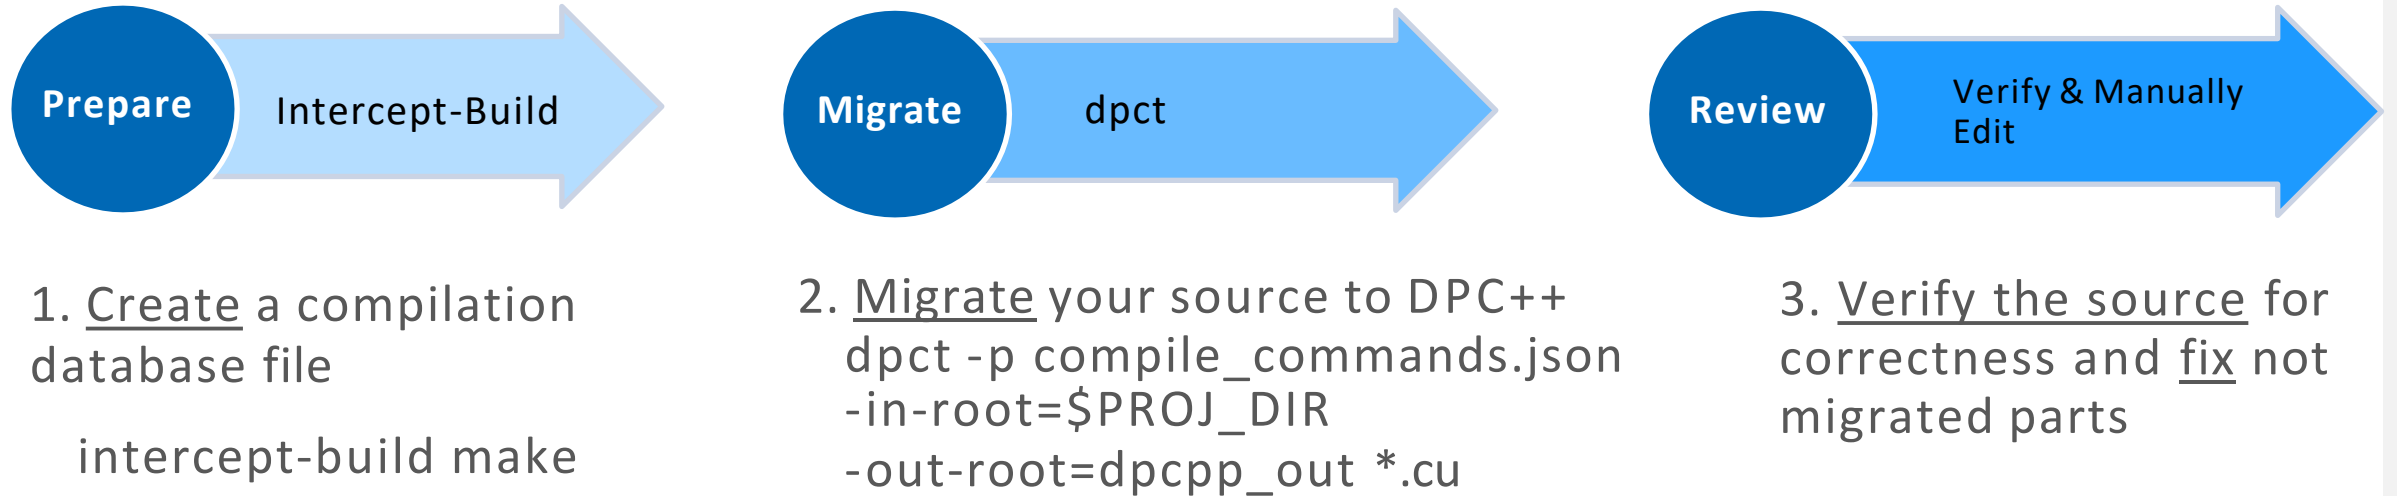
Inline comments are provided to help developers finish porting the application

Intel® DPC++ Compatibility Tool is Intel's implementation, available in the Base Toolkit

¹Intel estimates as of September 2021. Based on measurements on a set of 70 HPC benchmarks and samples, with examples like Rodinia, SHOC, PENNANT. Results may vary.

Intel® DPC++ Compatibility Tool

Migration of Large Code Bases



BlackScholes Demo

SYCLOMATIC on BlackScholes NVIDIA CUDA example

- `git clone https://github.com/NVIDIA/cuda-samples.git`
- `cd Samples/5_Domain_Specific/BlackScholes`
- `make clean`
- `intercept-build make`
- `mkdir ../migration`
- `dpct --cuda-include-path /usr/local/cuda-12/include \`
`-p compile_commands.json --in-root=. --out-root=../migration`
- `cp Makefile ../migration`

Changes to make command line

- `make: *** No rule to make target 'BlackScholes.cu', needed by 'BlackScholes.o'. Stop.`

```
BlackScholes.o:BlackScholes.cu =>  
BlackScholes.o:BlackScholes.dp.cpp
```

- `BlackScholes.dp.cpp:34:10: fatal error: sycl/sycl.hpp: No such file or directory`

```
make "NVCC=icpx -fsycl"
```

- `icpx: error: unsupported option '--threads'; did you mean '-mthreads'?`
`icpx: error: unknown argument: '-maxrregcount=16'`
`icpx: error: unknown argument: '-gencode'`
`make "NVCC=icpx -fsycl" ALL_CCFLAGS="" GENCODE_FLAGS="-g -O3"`

Changes to BlackScholes.dp.cpp

- BlackScholes.dp.cpp:106:3: error: use of undeclared identifier 'findCudaDevice'
checkCudaErrors(DPCT_CHECK_ERROR(
findCudaDevice(argc, (const char **)argv);
getLastCudaError("BlackScholesGPU() execution failed\n");
- add to BlackScholes.dp.cpp
`void checkCudaErrors(int err){}`
`void findCudaDevice(int argc, const char** argv){}`
`void getLastCudaError(char * strErr){}`

Different Performance

CUDA

- Executing Black-Scholes GPU kernel (512 iterations)...
- Options count : 8000000
- BlackScholesGPU() time : 0.158816 msec
- Effective memory bandwidth: 503.726277 GB/s
- Gigaoptions per second : **50.372628**
- BlackScholes, Throughput = 50.3726 GOptions/s, Time = 0.00016 s, Size = 8000000 options, NumDevsUsed = 1, Workgroup = 128

SYCL

- Executing Black-Scholes GPU kernel (512 iterations)...
- Options count : 8000000
- BlackScholesGPU() time : 1.463035 msec
- Effective memory bandwidth: 54.680848 GB/s
- Gigaoptions per second : **5.468085**
- BlackScholes, Throughput = 5.4681 GOptions/s, Time = 0.00146 s, Size = 8000000 options, NumDevsUsed = 1, Workgroup = 128

Checking app is running on correct GPU

- Check what accelerators are available

```
[opencl:acc:0] Intel(R) FPGA Emulation Platform for OpenCL(TM), Intel(R) FPGA Emulation Device 1.2 [2023.16.6.0.22_223734]  
[opencl:cpu:1] Intel(R) OpenCL, 13th Gen Intel(R) Core(TM) i9-13900HX 3.0 [2023.16.6.0.22_223734]  
[opencl:gpu:2] Intel(R) OpenCL HD Graphics, Intel(R) Graphics [0xa788] 3.0 [22.24.23453]  
[ext_oneapi_level_zero:gpu:0] Intel(R) Level-Zero, Intel(R) Graphics [0xa788] 1.3 [1.3.23453]  
[ext_oneapi_cuda:gpu:0] NVIDIA CUDA BACKEND, NVIDIA GeForce RTX 4090 Laptop GPU 8.8 [CUDA 12.0]
```

- Run app with SYCL_PI_TRACE

```
SYCL_PI_TRACE=1 ./Blackscholes  
SYCL_PI_TRACE[all]: Selected device: -> final score = 100  
SYCL_PI_TRACE[all]: platform: Intel(R) Level-Zero  
SYCL_PI_TRACE[all]: device: Intel(R) Graphics [0xa788]  
[./BlackScholes] - Starting...
```

Running on Intel GPU
not NVIDIA GPU !

Force app to use NVIDIA GPU using SYCL_DEVICE_FILTER

- SYCL_DEVICE_FILTER=cuda SYCL_PI_TRACE=1 ./BlackScholes

```
SYCL_PI_TRACE[basic]: Plugin found and successfully loaded: libpi_cuda.so [ PluginVersion: 12.27.1 ]
SYCL_PI_TRACE[all]: Requested device_type: info::device_type::automatic
SYCL_PI_TRACE[all]: Selected device: -> final score = 500
SYCL_PI_TRACE[all]:   platform: NVIDIA CUDA BACKEND
SYCL_PI_TRACE[all]:   device: NVIDIA GeForce RTX 4090 Laptop GPU
[./BlackScholes] - Starting...
Initializing data...
...allocating CPU memory for options.
...allocating GPU memory for options.
...generating input data in CPU mem.
...copying input data to GPU mem.
Data init done.

Executing Black-Scholes GPU kernel (512 iterations)...
terminate called after throwing an instance of 'sycl::_V1::runtime_error'
  what():  Native API failed. Native API returns: -42 (PI_ERROR_INVALID_BINARY) -42 (PI_ERROR_INVALID_BINARY)
Aborted
```

Invalid Binary: looks like wrong options were used in initial build!

Rebuilding and Running fixes the problem

- App should have been built with 'option' -fsycl-targets=nvptx64-nvidia-cuda
- make NVCC="icpx -fsycl" ALL_CCFLAGS="" GENCODE_FLAGS="-g -O3 -fsycl-targets=nvptx64-nvidia-cuda "

```
SYCL_PI_TRACE[all]: Selected device: -> final score = 1500
SYCL_PI_TRACE[all]: platform: NVIDIA CUDA BACKEND
SYCL_PI_TRACE[all]: device: NVIDIA GeForce RTX 4090 Laptop
[./BlackScholes] - Starting...
Initializing data...
...allocating CPU memory for options
...allocating GPU memory for options
...generating input data in CPU memory
...copying input data to GPU memory
Data init done.

Executing Black-Scholes GPU kernel (512 iterations)...
Options count          : 8000000
BlackScholesGPU() time  : 0.162633 msec
Effective memory bandwidth: 491.905667 GB/s
Gigaoptions per second  : 49.190567


BlackScholes, Throughput = 49.1906 GOptions/s, Time = 0.00016
DevsUsed = 1, Workgroup = 128
```

SYCL same Performance as CUDA

```
SYCL_PI_TRACE[basic]: Plugin found and successfully loaded: libpi_cuda.so [ PluginVersion: 12.27.1 ]
SYCL_PI_TRACE[all]: Requested device_type: info::device_type::automatic
SYCL_PI_TRACE[all]: Selected device: -> final score = 1500
SYCL_PI_TRACE[all]:   platform: NVIDIA CUDA BACKEND
SYCL_PI_TRACE[all]:   device: NVIDIA GeForce RTX 4090 Laptop GPU
[./BlackScholes] - Starting...
Initializing data...
...allocating CPU memory for options.
...allocating GPU memory for options.
...generating input data in CPU mem.
...copying input data to GPU mem.
Data init done.

Executing Black-Scholes GPU kernel (512 iterations)...
Options count          : 8000000
BlackScholesGPU() time  : 0.160705 msec
Effective memory bandwidth: 497.806309 GB/s
Gigaoptions per second  : 49.780631

BlackScholes, Throughput = 49.7806 GOptions/s, Time = 0.00016 s, Size = 8000000 options, NumDevsUsed = 1, Workgroup = 128
```



I
and
YOU


Backup

Lammps offloading

Stephen Blair-Chappell

Test Application

LAMMPS Molecular Dynamics Simulator



<https://www.lammps.org/>

Machine 1 Spec - Workstation

```
l58345@011-n001:~/ILDevCON/lammps-offload-OneAPI/LAB3/1-Model-Offload$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:             Little Endian
Address sizes:          46 bits physical, 48 bits virtual
CPU(s):                 24
On-line CPU(s) list:    0-23
Thread(s) per core:     2
Core(s) per socket:     12
Socket(s):              1
NUMA node(s):           1
Vendor ID:              GenuineIntel
CPU family:              6
Model:                  85
Model name:             Intel(R) Core(TM) i9-10920X CPU @ 3.50GHz
Stepping:               7
CPU MHz:                1200.315
CPU max MHz:            4800.0000
CPU min MHz:            1200.0000
BogoMIPS:               6999.82
Virtualization:         VT-x
L1d cache:              384 KiB
L1i cache:              384 KiB
L2 cache:               12 MiB
L3 cache:               19.3 MiB
NUMA node0 CPU(s):      0-23
Vulnerability Itlb multihit: KVM: Mitigation: Split huge pages
Vulnerability L1tf:       Not affected
Vulnerability Mds:        Not affected
Vulnerability Meltdown:   Not affected
Vulnerability Spec store bypass: Mitigation; Speculative Store Bypass disabled via prctl and seccomp
Vulnerability Spectre v1:  Mitigation; usercopy/swapgs barriers and __user pointer sanitization
Vulnerability Spectre v2:  Mitigation; Enhanced IBRS, IBPB conditional, RSB filling
Vulnerability Srbds:       Not affected
Vulnerability Tsx async abort: Mitigation; TSX disabled
Flags:                    fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush dts acp
                          i mmx fxsr sse sse2 ss ht tm pbe syscall nx pdpe1gb rdtscp lm constant_tsc art arch_per
                          fmon pebs bts rep_good nopl xtopology nonstop tsc cpuid aperfmperf pni pclmulqdq dtes64
                          monitor ds_cpl vmx est tm2 ssse3 sdbg fma cx16 xtpr pdcm pcid dca sse4_1 sse4_2 x2apic
                          movbe popcnt tsc_deadline_timer aes xsave avx f16c rdrand lahf_lm abm 3dnowprefetch cp
                          uid_fault epb cat_l3 cdp_l3 invpcid_single ssbd mba ibrs ibpb stibp ibrs_enhanced tpr_s
                          hadow vnmi flexpriority ept vpid ept_ad fsgsbase tsc_adjust bmi1 avx2 smep bmi2 erms in
                          vpcid cqm mpx rdt_a avx512f avx512dq rdseed adx smap clflushopt clwb intel_pt avx512cd
                          avx512bw avx512vl xsaveopt xsavec xgetbv1 xsaves cqm_llc cqm_occup_llc cqm_mbm_total cq
                          m_mbm_local dtherm ida arat pln pts hwp hwp_act_window hwp_epp hwp_pkg_req avx512_vnni
                          md clear flush_l1d arch_capabilities
```

Intel(R) Core(TM) i9-10920X CPU @ 3.50GHz

| | |
|----------------------|------|
| CPU(s): | 24 |
| On-line CPU(s) list: | 0-23 |
| Thread(s) per core: | 2 |

avx2 smep bmi2 erms in
wb intel_pt avx512cd
_llc cqm_mbm_total cq
_pkg_req avx512_vnni


Machine 2 Spec - Laptop

```
stephen@stephen-Swift-SF314-510G:~/dv/OneAPI-Package-1/LAMMPS$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
Address sizes:         39 bits physical, 48 bits virtual
CPU(s):                8
On-line CPU(s) list:   0-7
Thread(s) per core:    2
Core(s) per socket:    4
Socket(s):             1
NUMA node(s):         1
Vendor ID:             GenuineIntel
CPU family:            6
Model:                 140
Model name:            11th Gen Intel(R) Core(TM) i5-1135G7 @ 2.40GHz
```

```
u58345@s011-n001:~/ILDevCON/lammps-offload-OneAPI/LAB3/1-Model-Offload$ lspci | grep VGA
1c:00.0 VGA compatible controller: Intel Corporation Device 4905 (rev 01)
6a:00.0 VGA compatible controller: Intel Corporation Device 4905 (rev 01)
```

Graphics processor table


| PCI IDs | Name | Architecture | Codename |
|---------|------------------------------|--------------|----------|
| 4905 | Intel® Iris® Xe MAX Graphics | Xe | DG1 |



Analyze



Goal of Analysis



1. Find the 'hot spots'



2. if possible, predict benefit of offloading

Three Tools

APS – (Application
Performance Snapshot)

Advisor - model
offloading

VTune - Profiler

APS – (Application Performance Snapshot)

Application Performance Snapshot

Application: *Imp*
Report creation date: 2021-09-23 06:58:02
Number of ranks: 16
Ranks per node: 16
HW Platform: *Intel(R) microarchitecture code named Tigerlake*
Frequency: 2.42 GHz
Logical Core Count per node: 8
Collector type: *Event-based sampling driver, Event-based counting driver*

7.87 s CPU
Elapsed Time

1 IPC Rate

0 SP GFLOPS

1.43 DP GFLOPS

2.88 GHz Average CPU Frequency

GPU Utilization when Busy

3.6%

| EU State | % of EUs |
|----------|----------|
| Active | 3.6% |
| Idle | 83% |
| Stalled | 13.4% |

GPU Occupancy
9.1% of Peak Value

Memory Footprint

Resident
113.19 MB

Resident per Node
1811 MB

Virtual
2646.12 MB

Virtual Per Node
42338 MB

MPI Time

4.3 s

54.64% of Elapsed Time

MPI Imbalance
2.38 s
30.28% of Elapsed Time

| TOP 5 MPI Functions | % of Elapsed Time |
|---------------------|-------------------|
| MPI_Bcast | 28.84% |
| MPI_Wait | 9.86% |
| MPI_Allreduce | 7.87% |
| MPI_Init | 5.4% |
| MPI_Sendrecv | 1.69% |

Memory Stalls

7.2% of Pipeline Slots

Cache Stalls
19.2% of Cycles

DRAM Stalls
7.5% of Cycles

| DRAM Bandwidth | |
|----------------|----------|
| Average | 9 GB/s |
| Peak | 8.9 GB/s |
| Bound | 0% |

Vectorization

0%

Instruction Mix

SP FLOPs
0% of uOps

DP FLOPs
6.9% of uOps

Packed:
0% from DP FP
128-bit: 0%
256-bit: 0%
512-bit: 0%

Scalar:
100% from DP FP

Non-FP
93.1% of uOps

FP Arith/Mem Rd Instr. Ratio
0.21

FP Arith/Mem Wr Instr. Ratio
0.73

Your application may underutilize the GPU.

Run a [GPU Offload \(Preview\)](#) or a [GPU Compute/Media Hotspots \(Preview\)](#) analysis with VTune Profiler to discover how to better utilize the GPU.

| | Current run | Target | Tuning Potential |
|---------------------------|-------------|--------|------------------|
| MPI Time | 54.64% | <10% | |
| Memory Stalls | 7.2% | <20% | |
| Vectorization | 0% | >70% | |
| GPU Utilization when Busy | 3.6% | >80% | |

Advisor - model offloading

Top Metrics

5.600x

Speed-up for Accelerated Code

1.661x

Amdahl's Law Speed Up

48%

Fraction of Accelerated Code

20


Number of Offloads

Program Metrics



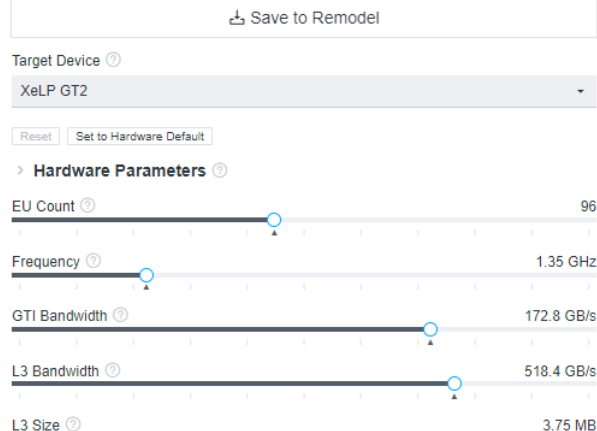
| | | | |
|---|---------|-------------------------------|--------------------------------|
| Program Time on Host After Acceleration | 2.088s | Speed-up for Accelerated Code | 5.600x |
| Time on Target | 85.9ms | Amdahl's Law Speed Up | 1.661x |
| Time in MPI calls | 88.0ms | Fraction of Accelerated Code | 48% |
| Non-Accelerable Time | 254.0ms | Number of Offloads | 20 |
| Data Transfer Tax | 0s | CPU Threads | 1 |
| Kernel Launch Tax | 10.4ms | Target Platform | XeLP GT2 |
| | | Baseline Platform | Intel(R) Xeon(R) Gold 6128 CPU |

Offload Bounded By



| | |
|---------------|-----|
| Compute | 11% |
| L3 Cache BW | 5% |
| GTI BW | 0% |
| LLC BW | 0% |
| Memory BW | 24% |
| Atomics | 0% |
| Latencies | 0% |
| Data Transfer | 0% |
| Launch Tax | 9% |
| Dependency | 0% |
| Trip Count | 0% |
| Unknown | 0% |
| Non Offloaded | 52% |

Modeling Parameters



Top Offloaded

| Loop/Function | Execution Time | Speed-Up | Bounded By | Data Transfer |
|---|---------------------------|----------|------------|---------------|
| [loop in LAMMPS_NS::PairLJCutIntel::eval<(int)0,(int)0,(int)1,float,double>Somp\$parallel@196 at pair_lj_cut_intel.cpp:219] | CPU 408.0ms GPU 40.3ms | 10.117x | Compute | 17.7MB |
| [loop in LAMMPS_NS::Atom::sort at atom.cpp:2086] | CPU 304.1ms GPU 7.1ms | 42.653x | DRAM BW | 18.1MB |
| [loop in LAMMPS_NS::FixLangevin::post_force_templated<(int)0,(int)0,(int)0,(int)0,(int)0> at fix_langevin.cpp:643] | CPU 259.7ms GPU 8.1ms | 32.039x | DRAM BW | 9.33MB |
| [loop in LAMMPS_NS::BondFENEIntel::eval<(int)0,(int)0,(int)1,float,double>Somp\$parallel@132 at bond_fene_intel.cpp:162] | CPU 123.9ms GPU 4.6ms | 27.103x | L3 BW | 9.21MB |
| [loop in LAMMPS_NS::FixNVEIntel::initial_integrate at fix_nve_intel.cpp:78] | CPU 112.2ms GPU 4.7ms | 23.654x | DRAM BW | 9.27MB |

Top Non-Offloaded

| Loop/Function | Execution Time | Bounded By | Why Not Offloaded | Data Transfer |
|--|---------------------------|--------------------------------|---|---------------|
| [loop in fi_getinfo] | CPU 1.3s GPU 1.3s | Launch Tax, Latencies, DRAM BW | Not profitable: Launch Tax, Latencies, DRAM Bandwidth Time is greater than other execution time components on a Target Device | 258kB |
| [loop in LAMMPS_NS::Neighbor::build at neighbor.cpp:2328] | CPU 255.8ms GPU 15.44s | Trip Counts, Latencies | Not profitable: Trip Counts, Latencies | 13.3MB |
| [loop in LAMMPS_NS::Replicate::command at replicate.cpp:655] | CPU 20.0ms GPU 97.4ms | Trip Counts, Latencies | Not profitable: Trip Counts, Latencies | 11.9MB |
| [loop in LAMMPS_NS::Neighbor::build at neighbor.cpp:2328] | CPU 20.0ms GPU 578.5ms | Trip Counts, Latencies | Not profitable: Trip Counts, Latencies | 1.04MB |

Steps to Offload Projection with Advisor

1. Run a **Survey**: get a list of hotspots

```
advisor -collect survey ...
```

- Sampling
- Binary Static Analysis
- Compiler & debug info

2. Run a **Trip Count**: count loop iteration

```
advisor -collect=tripcounts -target-device=gen9_gt2 .
```

- Trip count
- Cache simulation

3. Perform a **dependency analysis** [optional for quick modelling]

```
advisor -collect dependencies . . .
```

- Check memory accesses
- Loop selection heuristic

4. **Model** the Performance

```
advisor -collect projection -no-assume-dependencies . . .
```

- Generate HTML report

Expensive Steps

Top Metrics

5.600x

Speed-up for Accelerated Code

1.661x

Amdahl's Law Speed Up


48%

Fraction of Accelerated Code

20

Number of Offloads


Program Metrics



- Program Time on Host After Acceleration
- Time on Target
- Time in MPI calls
- Non-Accelerable Time
- Data Transfer Tax
- Kernel Launch Tax

| | | |
|---------|-------------------------------|--------------------------------|
| 2.088s | Speed-up for Accelerated Code | 5.600x |
| 85.9ms | Amdahl's Law Speed Up | 1.661x |
| 88.0ms | Fraction of Accelerated Code | 48% |
| 254.0ms | Number of Offloads | 20 |
| 0s | CPU Threads | 1 |
| 10.4ms | Target Platform | XeLP GT2 |
| | Baseline Platform | Intel(R) Xeon(R) Gold 6128 CPU |

Offload Bounded By



Modeling Parameters


Save to Remodel

Target Device

XeLP GT2

Reset Set to Hardware Default

Hardware Parameters

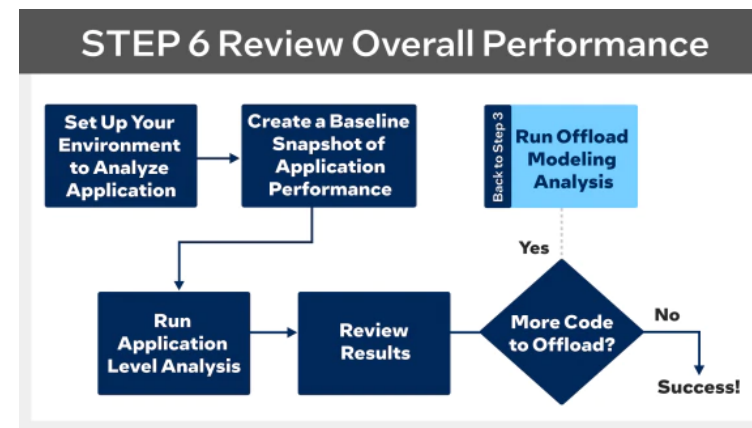
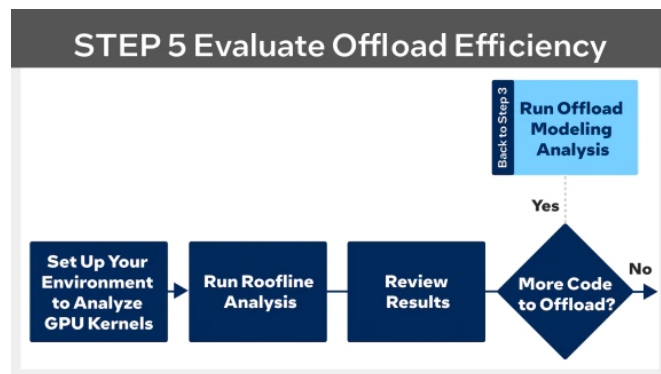
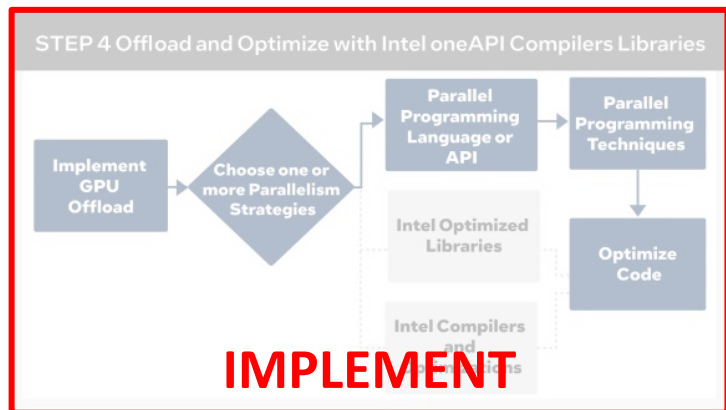
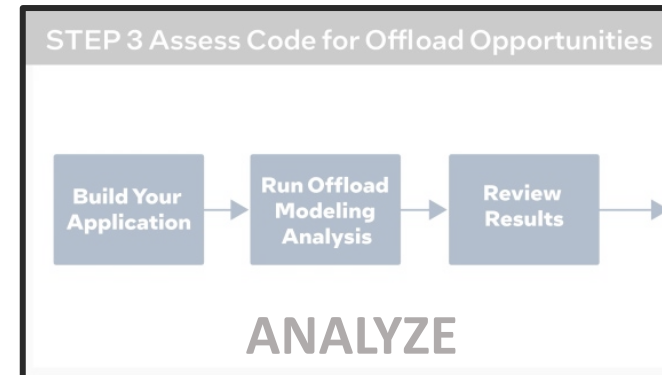
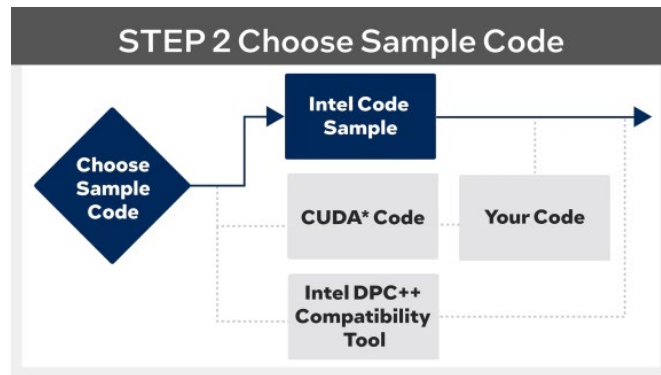
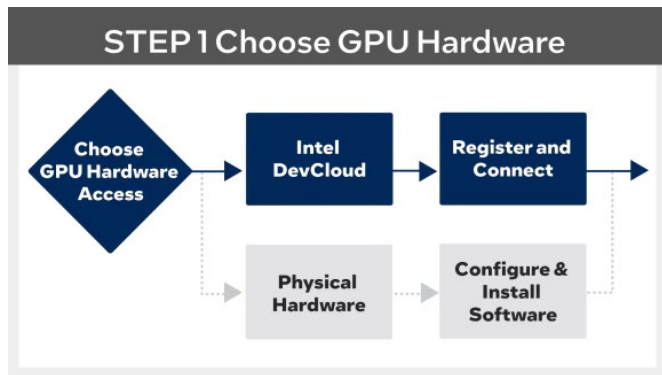


Top Offloaded

| Loop/Function | Execution Time | Speed-Up | Bounded By | Data Transfer |
|---|---------------------------|----------|------------|---------------|
| [loop in LAMMPS_NS::PairLJCutIntel::eval<(int)0,(int)0,(int)1,float,double>Somp\$parallel@196 at pair_lj_cut_intel.cpp:219] | CPU 408.0ms GPU 40.3ms | 10.117x | Compute | 17.7MB |
| [loop in LAMMPS_NS::Atom::sort at atom.cpp:2086] | CPU 304.1ms GPU 7.1ms | 42.653x | DRAM BW | 18.1MB |
| [loop in LAMMPS_NS::FixLangevin::post_force_templated<(int)0,(int)0,(int)0,(int)0,(int)0> at fix_langevin.cpp:643] | CPU 259.7ms GPU 8.1ms | 32.039x | DRAM BW | 9.33MB |
| [loop in LAMMPS_NS::BondFENEIntel::eval<(int)0,(int)0,(int)1,float,double>Somp\$parallel@132 at bond_fene_intel.cpp:162] | CPU 123.9ms GPU 4.6ms | 27.103x | L3 BW | 9.21MB |
| [loop in LAMMPS_NS::FixNVEIntel::initial_integrate at fix_nve_intel.cpp:78] | CPU 112.2ms GPU 4.7ms | 23.654x | DRAM BW | 9.27MB |

Top Non-Offloaded

| Loop/Function | Execution Time | Bounded By | Why Not Offloaded | Data Transfer |
|--|---------------------------|--------------------------------|---|---------------|
| [loop in fi_getinfo] | CPU 1.3s GPU 1.3s | Launch Tax, Latencies, DRAM BW | Not profitable: Launch Tax, Latencies, DRAM Bandwidth Time is greater than other execution time components on a Target Device | 258kB |
| [loop in LAMMPS_NS::Neighbor::build at neighbor.cpp:2328] | CPU 255.8ms GPU 15.44s | Trip Counts, Latencies | Not profitable: Trip Counts, Latencies | 13.3MB |
| [loop in LAMMPS_NS::Replicate::command at replicate.cpp:655] | CPU 20.0ms GPU 97.4ms | Trip Counts, Latencies | Not profitable: Trip Counts, Latencies | 11.9MB |
| [loop in LAMMPS_NS::Neighbor::build at neighbor.cpp:2328] | CPU 20.0ms GPU 578.5ms | Trip Counts, Latencies | Not profitable: Trip Counts, Latencies | 1.04MB |



<https://www.intel.com/content/www/us/en/developer/tools/oneapi/gpu-optimization-workflow.html>

TWO TYPES OF COMPILER

▪ **CLASSIC** [Same as was in Parallel Studio]

Deprecated and will be removed from product release in the second half of 2023

- icc
- icpc
- ifort

Offloading not supported

LLVM based [Totally new compilers]

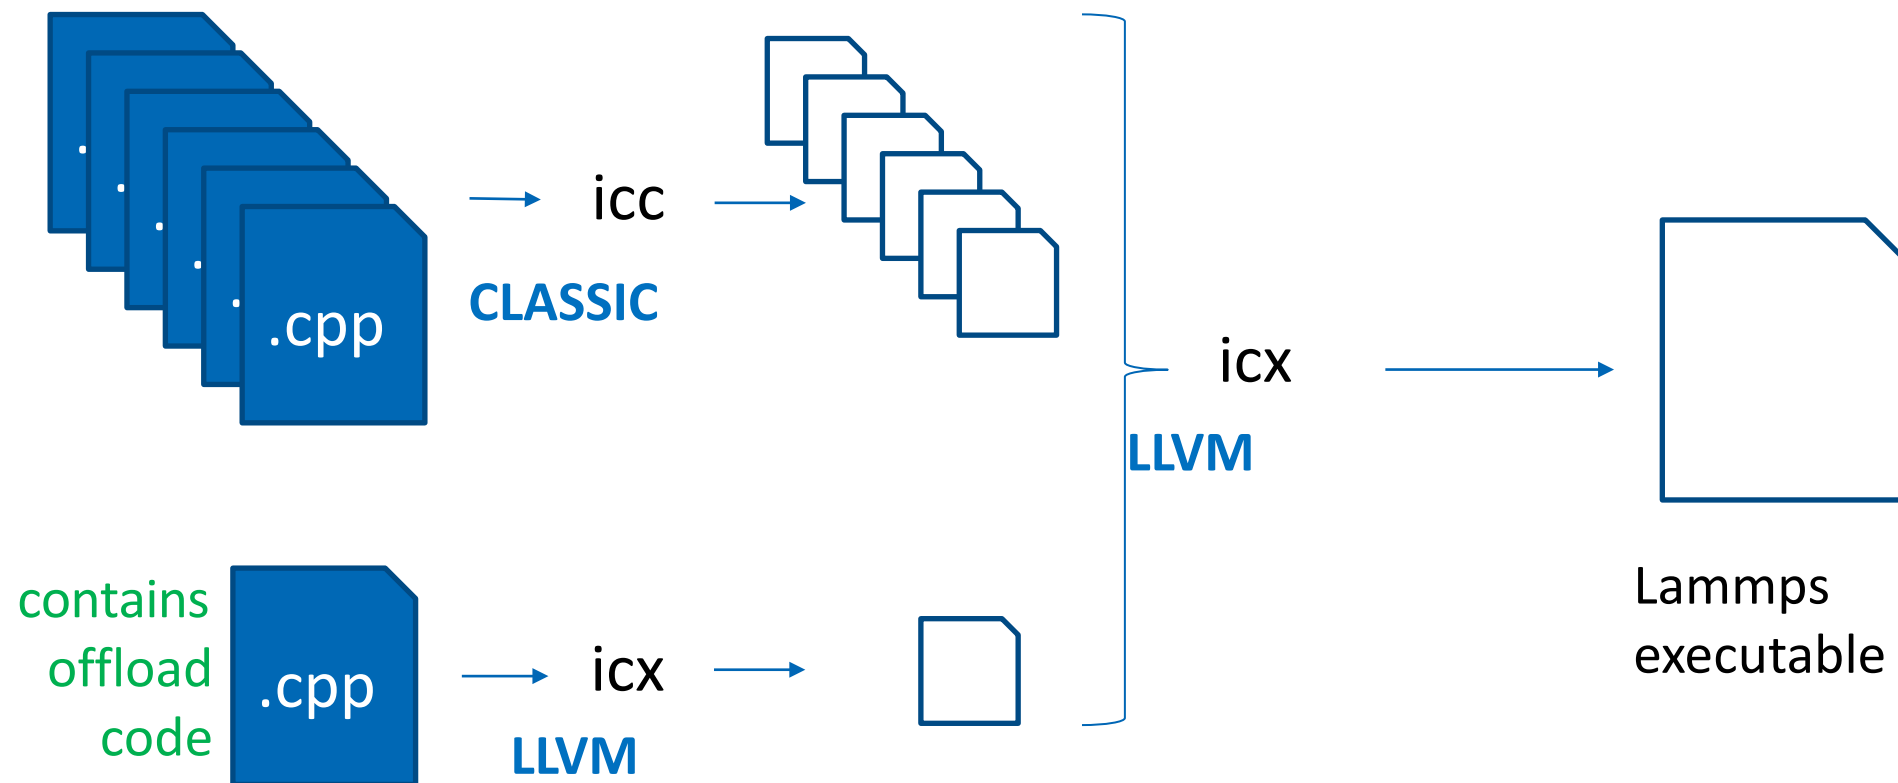
- icx
- dpcpp
- ifx

Offloading supported


All objects are binary compatible

One approach...

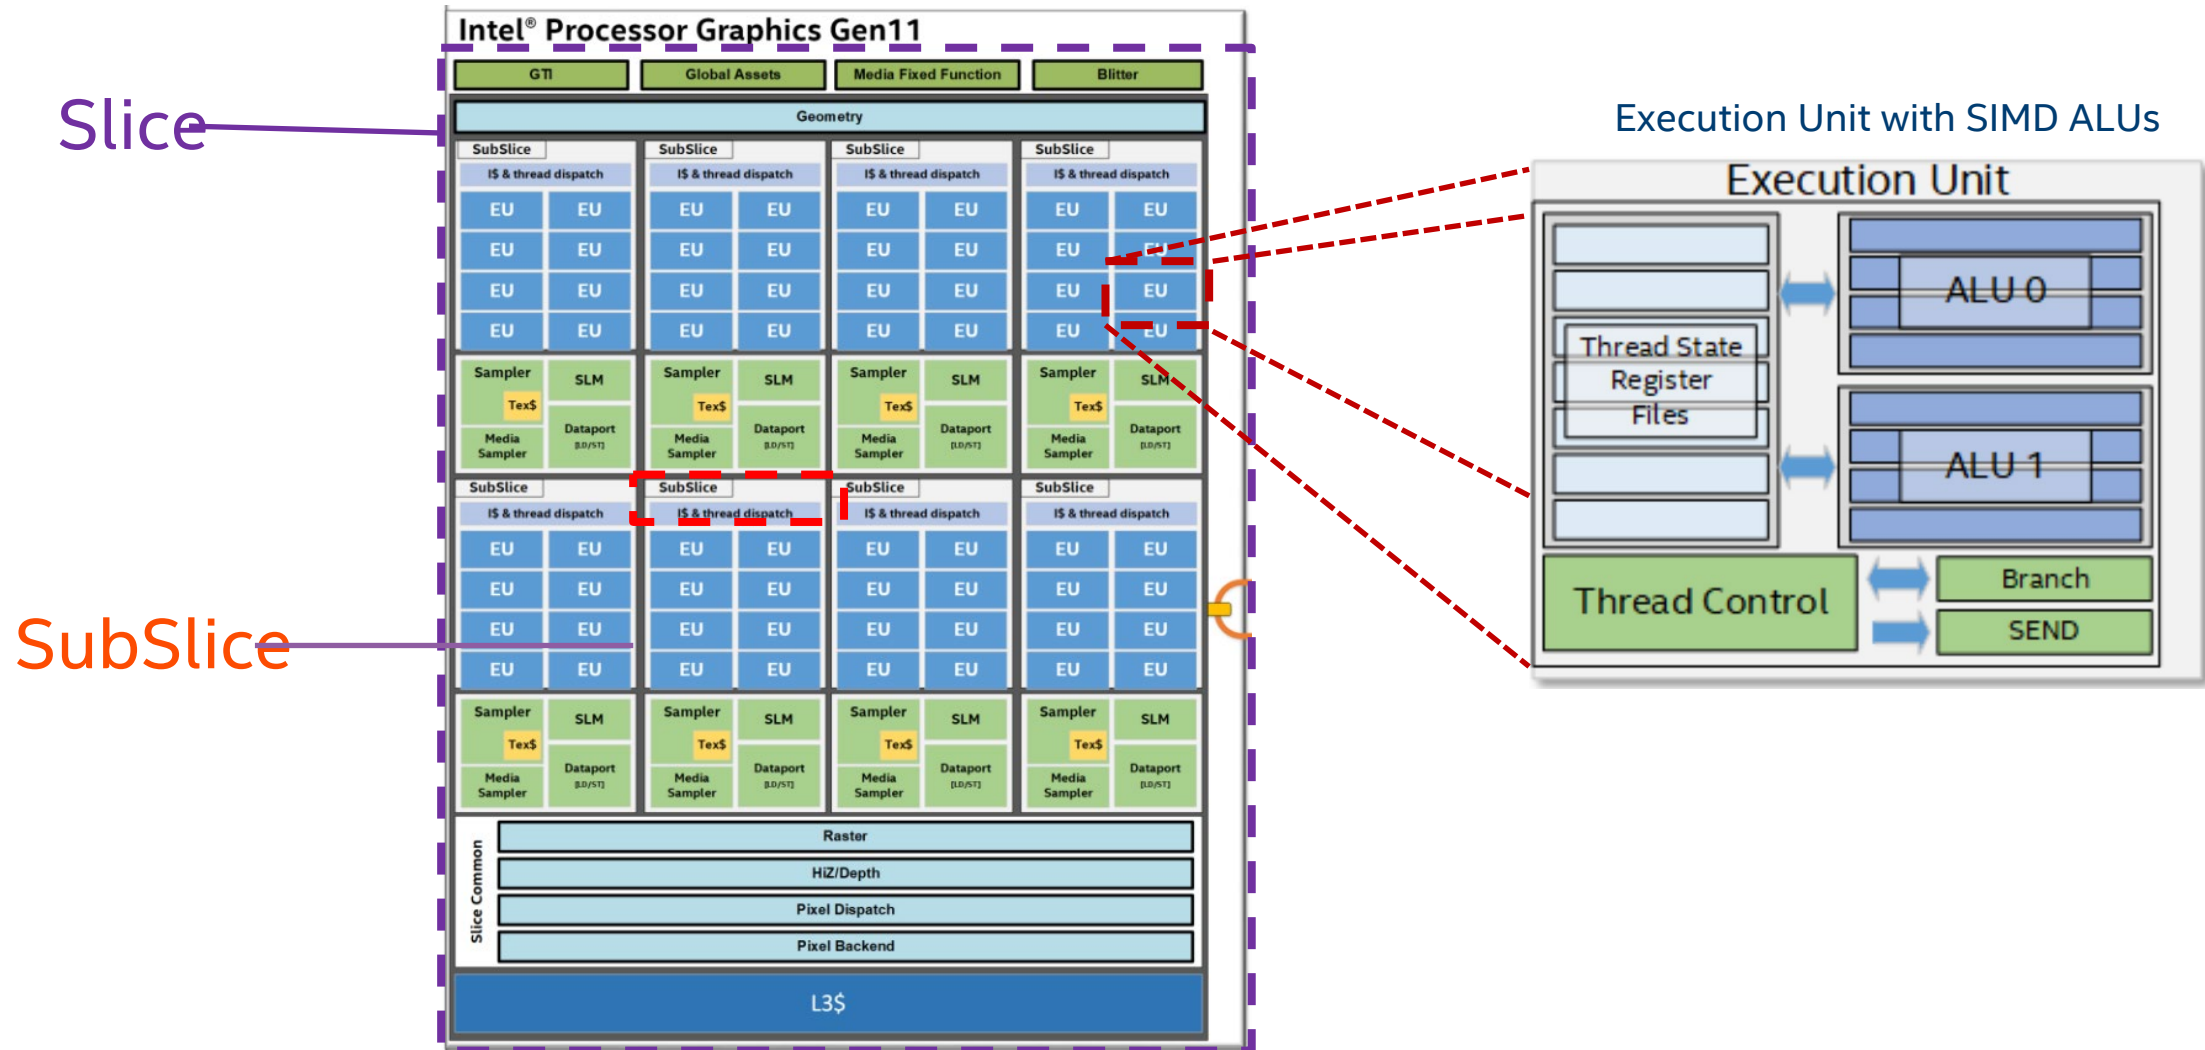
NB: icc deprecated mid 2023



... or alternatively



GPU Architecture



GPU Architecture



#pragma omp target

parallel for

Sends Code to target
but only on one sub-slice



OpenMP GPU Offload and OpenMP Constructs

- OpenMP GPU offload support all “normal” OpenMP constructs
 - E.g. parallel, for/do, barrier, sections, tasks, etc.
 - Not every construct will be useful
- Full threading model outside of a single GPU subslice **not supported**
 - No synchronization among subslices
 - No coherence and memory fence between among subslice L1 caches

GPU Architecture



#pragma omp target teams

parallel for

Share code across subslices
but only on one EU




GPU Architecture

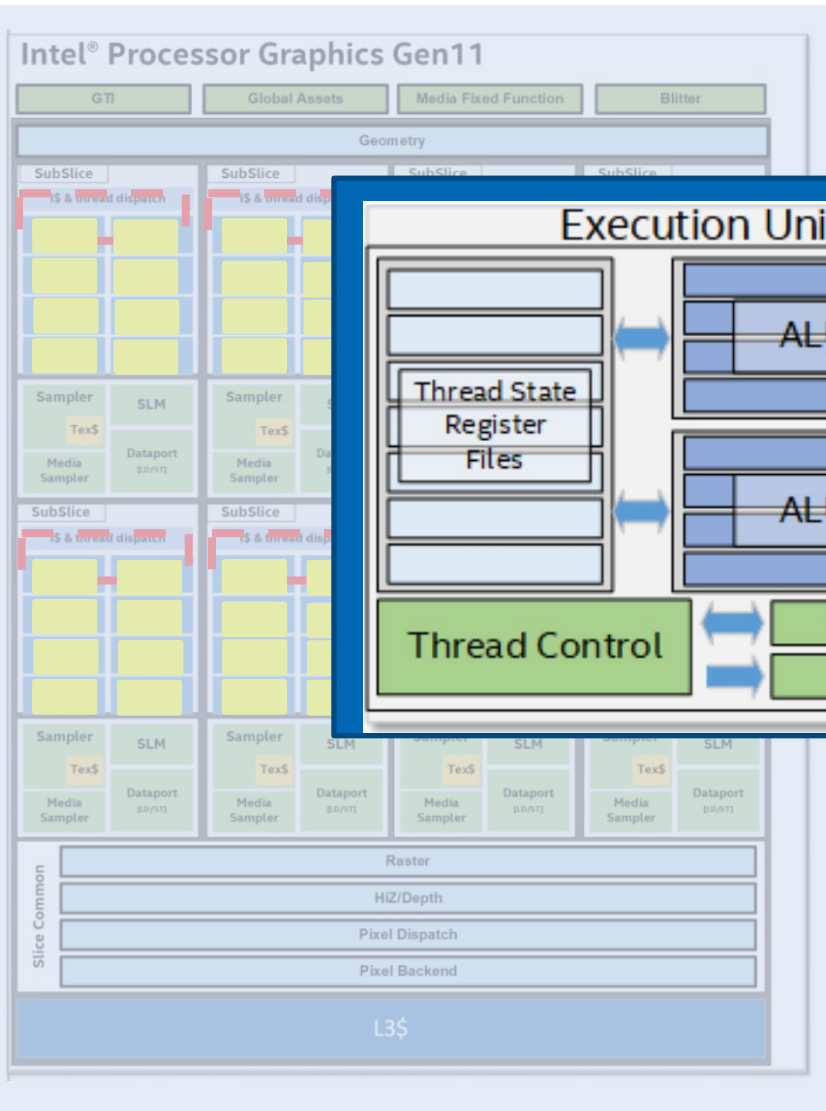


`#pragma omp target teams distribute parallel for`

Distributes loop iterations
across all EUs



GPU Architecture



#pragma omp target teams distribute simd parallel

Uses SIMD wide registers
in execution units



Experiment in Lammmps -Target

```
#if 1
    #pragma omp target map(to:ilist[iifrom:iito]) \
                        map(to:x[x_min:x_max]) \
                        map(tofrom:f[f_min:f_max]) \
                        map(from:ev_global[0:7]) \
                        map (tofrom:lj1,lj2,lj3,lj4,offset)
    #pragma omp teams distribute parallel for
#endif
// ----- END SBC code -----

for (int ii = iifrom; ii < iito; ii += iip) {
    const int i = ilist[ii];
    int itype, ptr_off;
    const FC_PACKED1_T * _noalias ljcl2oi;
    const FC_PACKED2_T * _noalias lj34i;
    if (!ONETYPE) {
        itype = x[i].w;
        ptr_off = itvne * ntvnes;
```

Experiment in Lammmps –Indirect Indexes

```
Start here *pair_lj_cut_intel.cpp SBC code
222 // ----- SBC code -----
223 #if 1
224 // find min and max of indirect indexes
225 int i_min=0;
226 int i_max=0;
227 int jlist_min=0;
228 int jlist_max=0;
229 int j_min=0;
230 int j_max=0;
231 int x_min=0;
232 int x_max=0;
233 int f_min=0;
234 int f_max=0;
235
236 for (int ii = iifrom; ii < iito; ii += iip) {
237     int i=0, itype=0, sbindex=0;
238
239     // get i_min and i_max
240     i = ilist[ii];
241     i_min=std::min(i,i_min);
242     i_max=std::max(i,i_max);
243
244     const int * _noalias const jlist = firstneigh[i];
245     int jnum = numneigh[i];
246
247     // we also need the j_min and j_max as this is used in x[j]
248     for (int jj = 0; jj < jnum; jj++) {
249         int j=0, jtype=0, sbindex=0;
250         if (!ONETYPE) {
251             sbindex = jlist[jj] >> SBBITS & 3;
252             j = jlist[jj] & NEIGHMASK;
253         } else
254             j = jlist[jj];
255
256         j_min=std::min(j,j_min);
257         j_max=std::max(j,j_max);
258     }
259     // ii
260     x_min = f_min = std::min(i_min,j_min);
261     x_max = f_max = std::max(i_max,j_max);
262
263     #if 0
264     printf("x_min: %d, x_max: %d ", x_min, x_max);
265     #endif
266     #endif
267
268     #if 1
269     #pragma omp target map(to:ilist[iifrom:iito]) \
270
```

```
// we also need the j_min and j_max as this
for (int jj = 0; jj < jnum; jj++) {
    int j=0, jtype=0, sbindex=0;
    if (!ONETYPE) {
        sbindex = jlist[jj] >> SBBITS & 3;
        j = jlist[jj] & NEIGHMASK;
    } else
        j = jlist[jj];

    j_min=std::min(j,j_min);
    j_max=std::max(j,j_max);
}
} // ii
x_min = f_min = std::min(i_min,j_min);
x_max = f_max = std::max(i_max,j_max);
```